Vector Stores vs Vector Databases

	Vector Store	Vector Database
What it is	A lightweight library or tool focused on storing and searching vectors efficiently	A full-featured database system designed for vector data at scale
Core Function	Simple similarity search - find the K nearest neighbors to a query vector	Advanced search with filters, metadata queries and database operations
Architecture	Usually runs in-memory or as a local file, single-machine operation	Distributed system with replication, sharding, and high availability
Best for	Prototypes, research, small applications with <1M vectors	Production systems, enterprise applications, billions of vectors

Quick Comparison

	Vector Store	Vector Database
Scale	~1M Vector	Billions +
Setup Time	Minutes	Hours / Days
Cost	Free / $	$$$ / $$$$
Query Speed	Microseconds	Milliseconds
Features	Basic Search	CRUD
Deployment	Local	Cloud
Examples	FAISS, Annoy, ChromaDB, ScaNN, NMSLIB	Pinecone, Weaviate, Qdrant, Milvus, Vespa

When to use

Vector Store	Vector Database
Building a proof of concept	Building production applications
Working with <1 million vectors	Need to scale beyond millions of vectors
Need fastest possible search speed	Require high availability and reliability
Have limited budget	Need advanced filtering and metadata search
Want full control over the implementation	Have multiple users / tenants
Building embedded applications	Want managed infrastructure