

Autonomous RAG

Autonomous RAG is a Retrieval Augmented Generation system where the LLM (or agent) is capable of reasoning, planning, acting, reflecting, and improving – on its own – without manual control over each step.

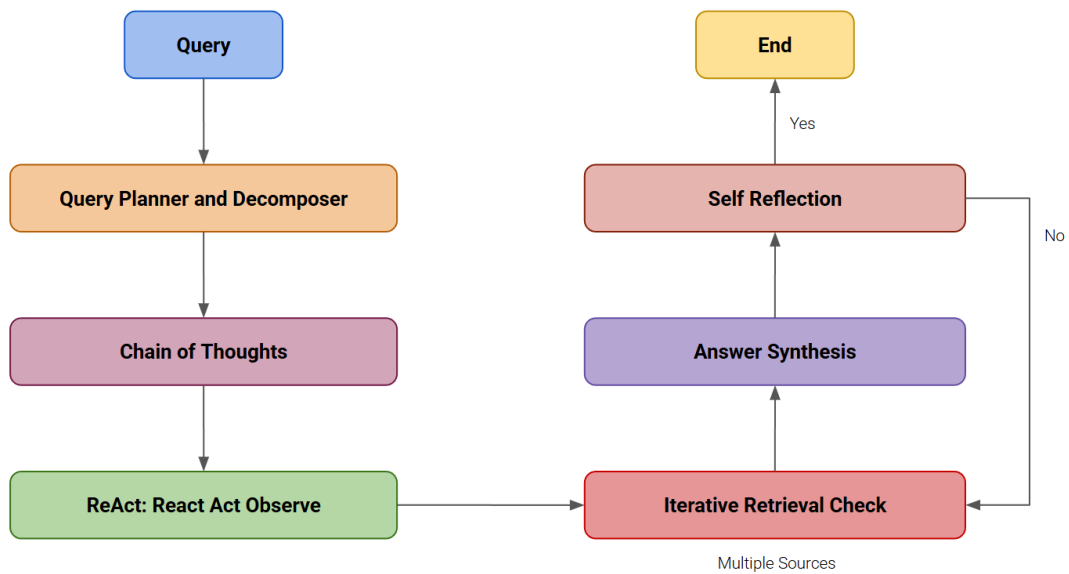
It combines

- Agentic reasoning (like ReAct or LangGraph agents)
- Self-reflection and self-correction
- Dynamic tool selection
- Multi-source retrieval

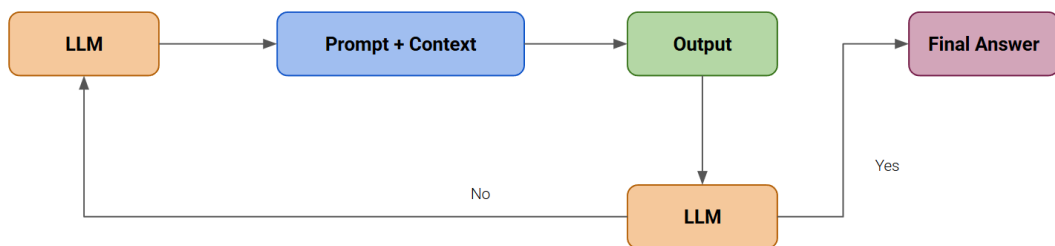
Core Components of Autonomous RAG

Component	Role
Planner Agent	Breaks complex queries into sub-questions
Tool Selector	Chooses between Wikipedia, ArXiv, vector DBs, APIs, etc
Retriever	Executes tool calls to retrieve relevant documents
Synthesizer	Uses LLM to generate the final answer
Reflector	Verifies whether context or answer is good enough
Retry Loop	Refines and retries if reflection fails
Memory (opt.)	Stores feedback, log bad queries, or improves prompts/tools

Complete Flow of Autonomous RAG

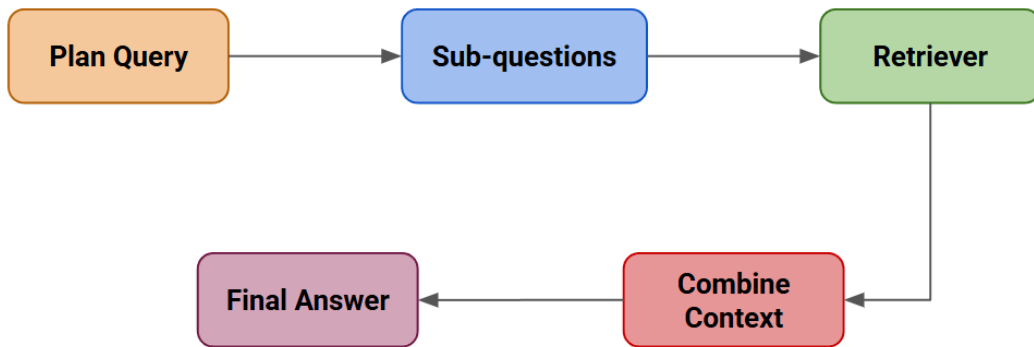


Self Reflection RAG

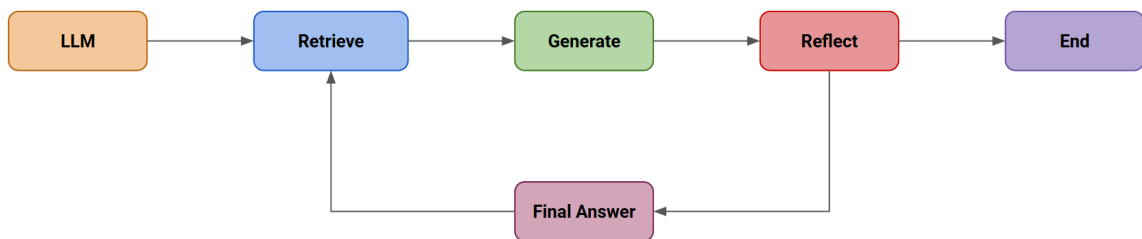


Query Planning and Decomposition RAG

	Chain of Thoughts	Query Planning and Decomposition
Purpose	Let the LLM reason step-by-step	Break a complex query into structured sub-queries
Style	Natural language reasoning path	Explicit sub-queries or formal question segments
Inspiration	Human-like scratchpad thinking	Structured task planning or modular Q&A
Agent Behavior	Think → Retrieve → Think → Answer	Plan all → Retrieve all → Answer once



Iterative Retrieval with Self-Reflection



Benefits

Feature	Advantage
Reflection	Reduces hallucination or overconfidence
Query refinement	Improves recall and relevance
Agentic behavior	Mimics human-like research process

Agentic RAG vs Autonomous RAG

Concept	Agentic RAG	Autonomous RAG
Definition	A RAG system that uses an agentic approach – where an LLM reasons, plans, and acts using tools	A RAG system that operates independently, with full self-management of planning, retrieving, reflection, and improvement
Focus	Structured reasoning and tool use (ReAct, LangGraph, etc)	Complete autonomy in task execution, retry, and learning
Behavior	Think → Act → Observe → Answer	Think → Act → Reflect → Retry → Learn → Answer
Retry Logic	Optional – usually static agent plans	Built-in retry/refine strategies (context + answer reflection)
Self-Reflection	May include it optionally	Core feature: reflects on retrieval and answers before finalizing
Tool use	Uses tool via agents (e.g., Wikipedia, SQL, ArXiv)	Selects and adapts tools dynamically based on reasoning
Planner	Often present (manual or LLM-generated plans)	Always present – triggers multi-step workflows adaptively
Learning Loop	Not always present	May log feedback, improve over time